



KUBERNETES

Container Management

Introduction and demonstration on real hardware, VMs.

Demo requirements: 3-4 nodes with CentOS 7.

Igor Dolzhikov, Openprovider

It all starts with one case



Several cases of controlled manually



A lot of cases without management




Controlled and balanced system



Container management systems

- Docker Swarm is native clustering for Docker.



Container management systems

- Apache Mesos - provides efficient resource isolation and sharing across distributed applications (used in Twitter and Apple for more than 10 000 containers).




Container management systems

- Kubernetes - a simple and flexible solution from Google, which thanks to the community is actively developing the last 2-3 years.



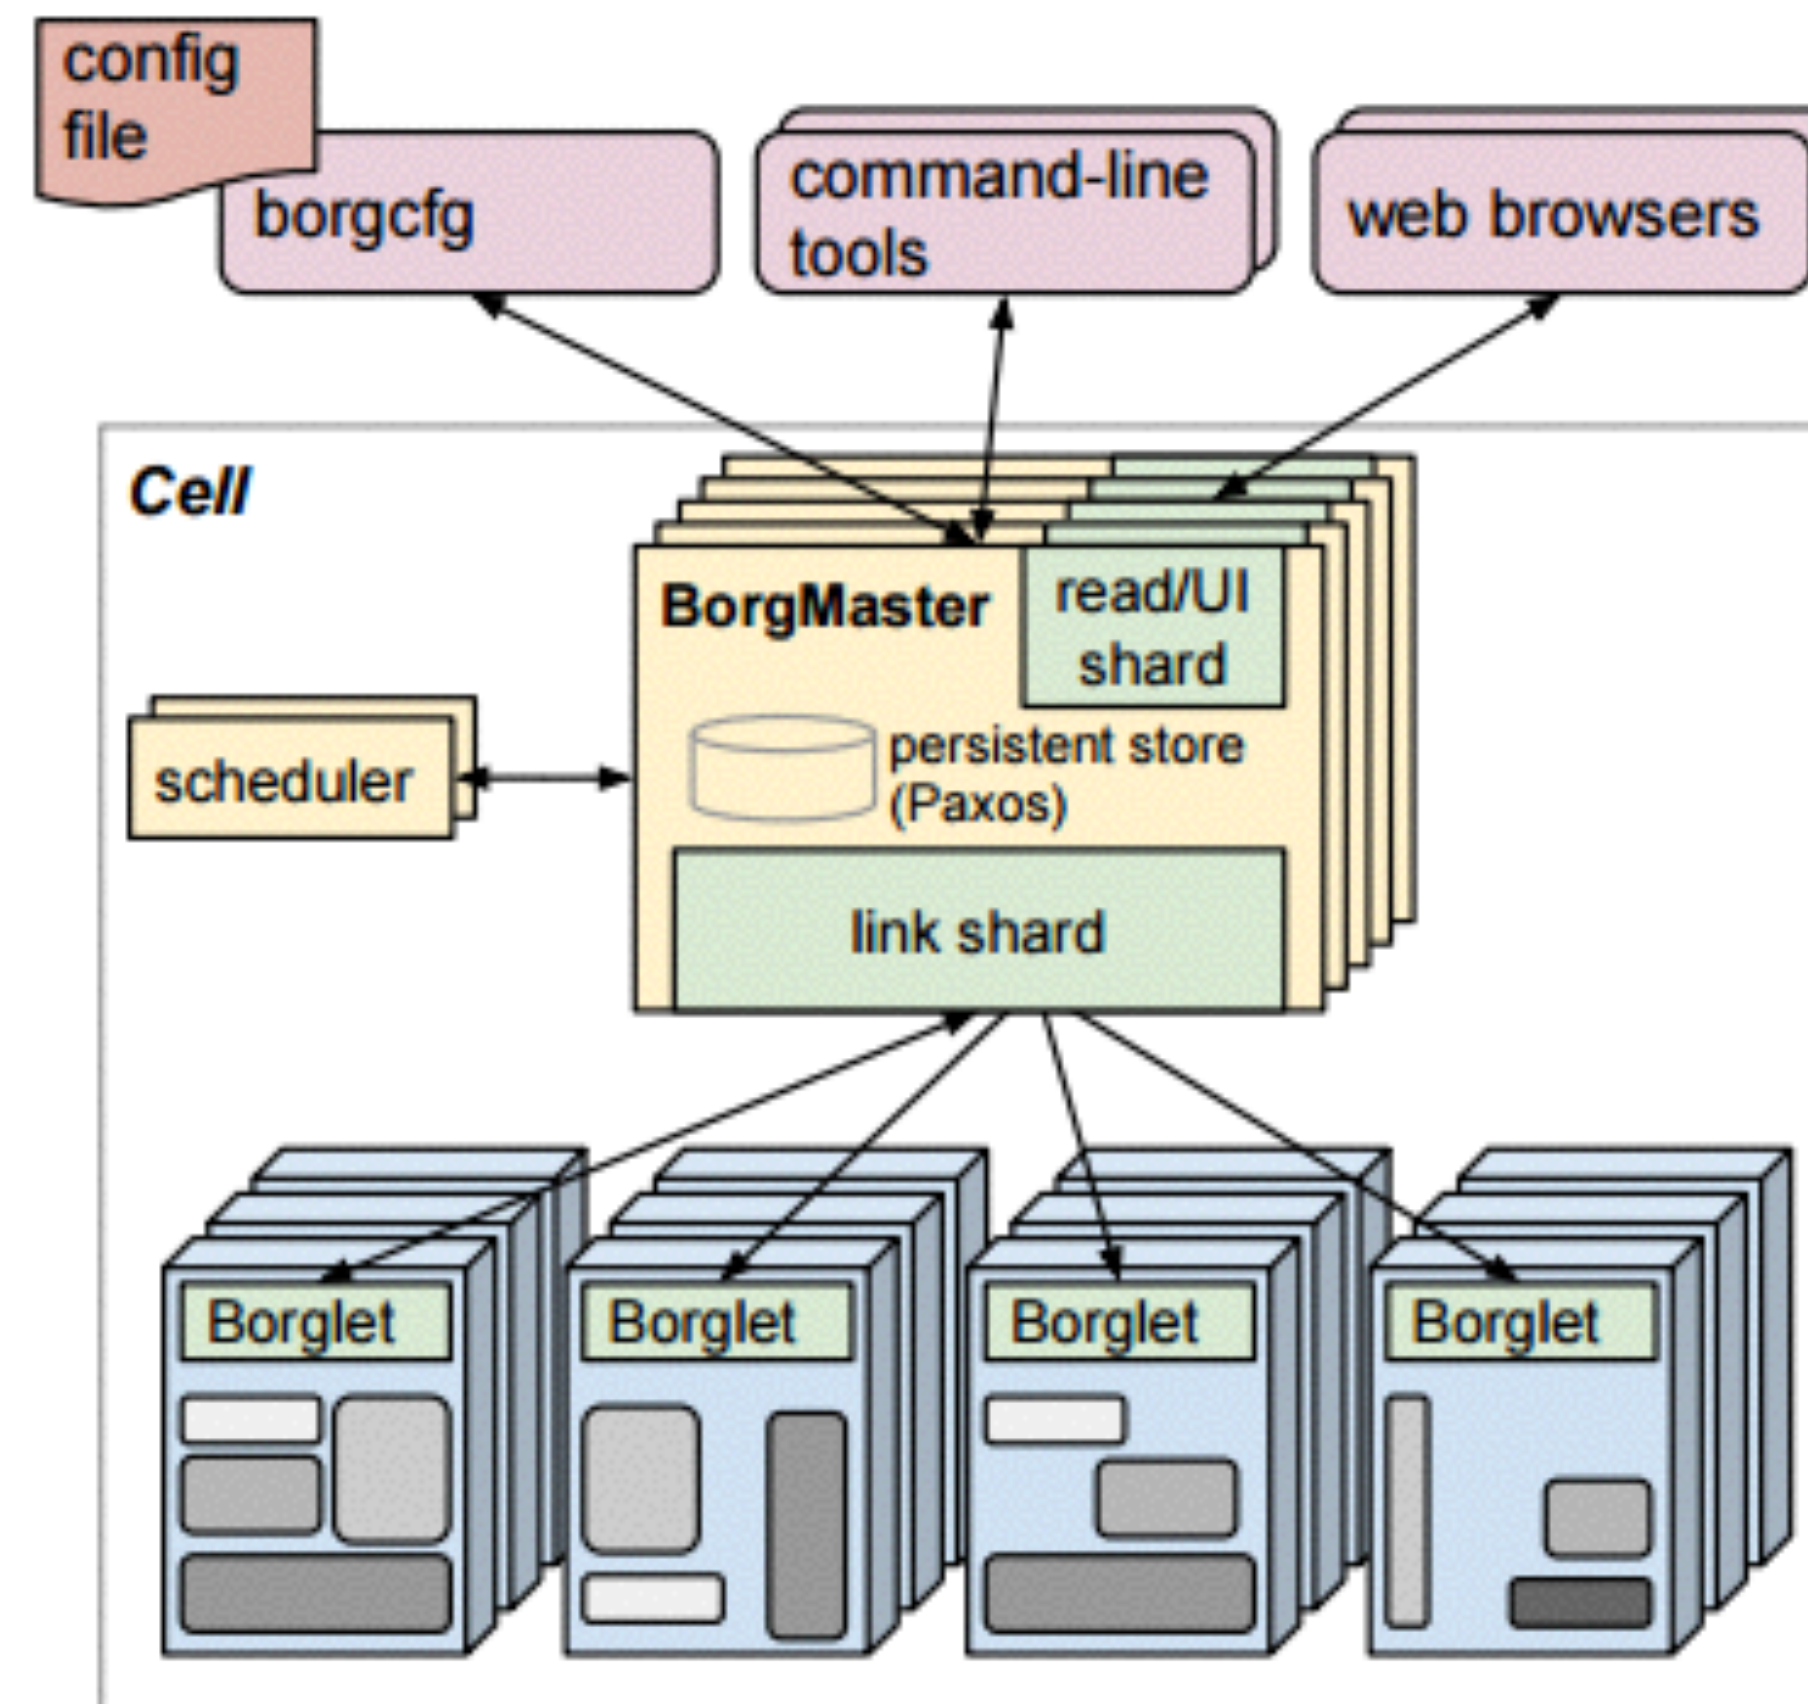
Container management systems

- And many others ...



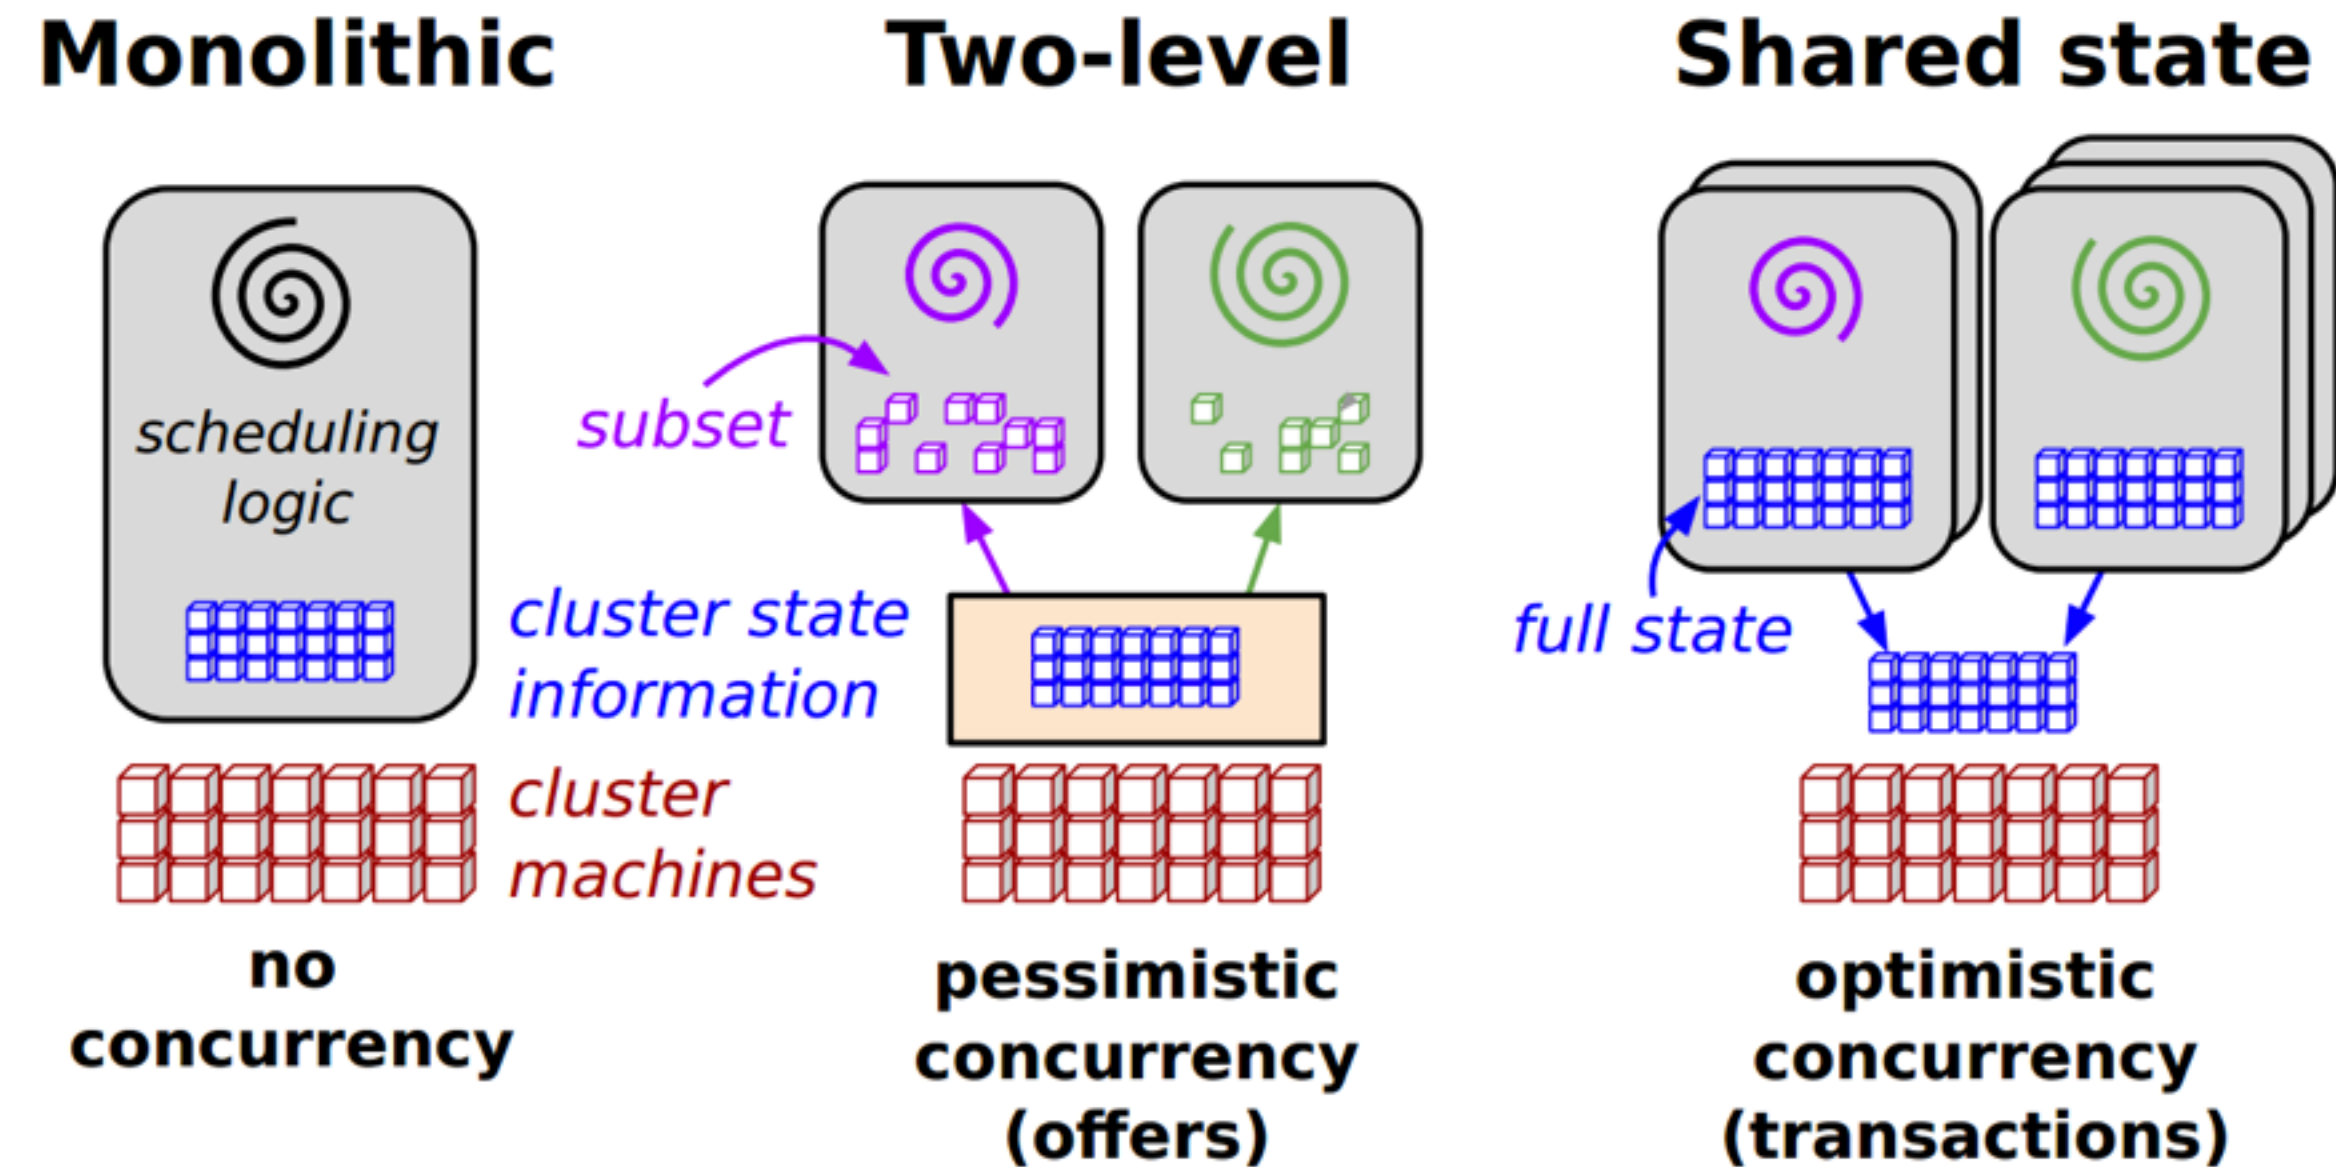
Kubernetes origin

- Borg - unified container-management system developed at Google.




Kubernetes origin

- Omega - an offspring of Borg, was driven by a desire to improve the software engineering of the Borg ecosystem.

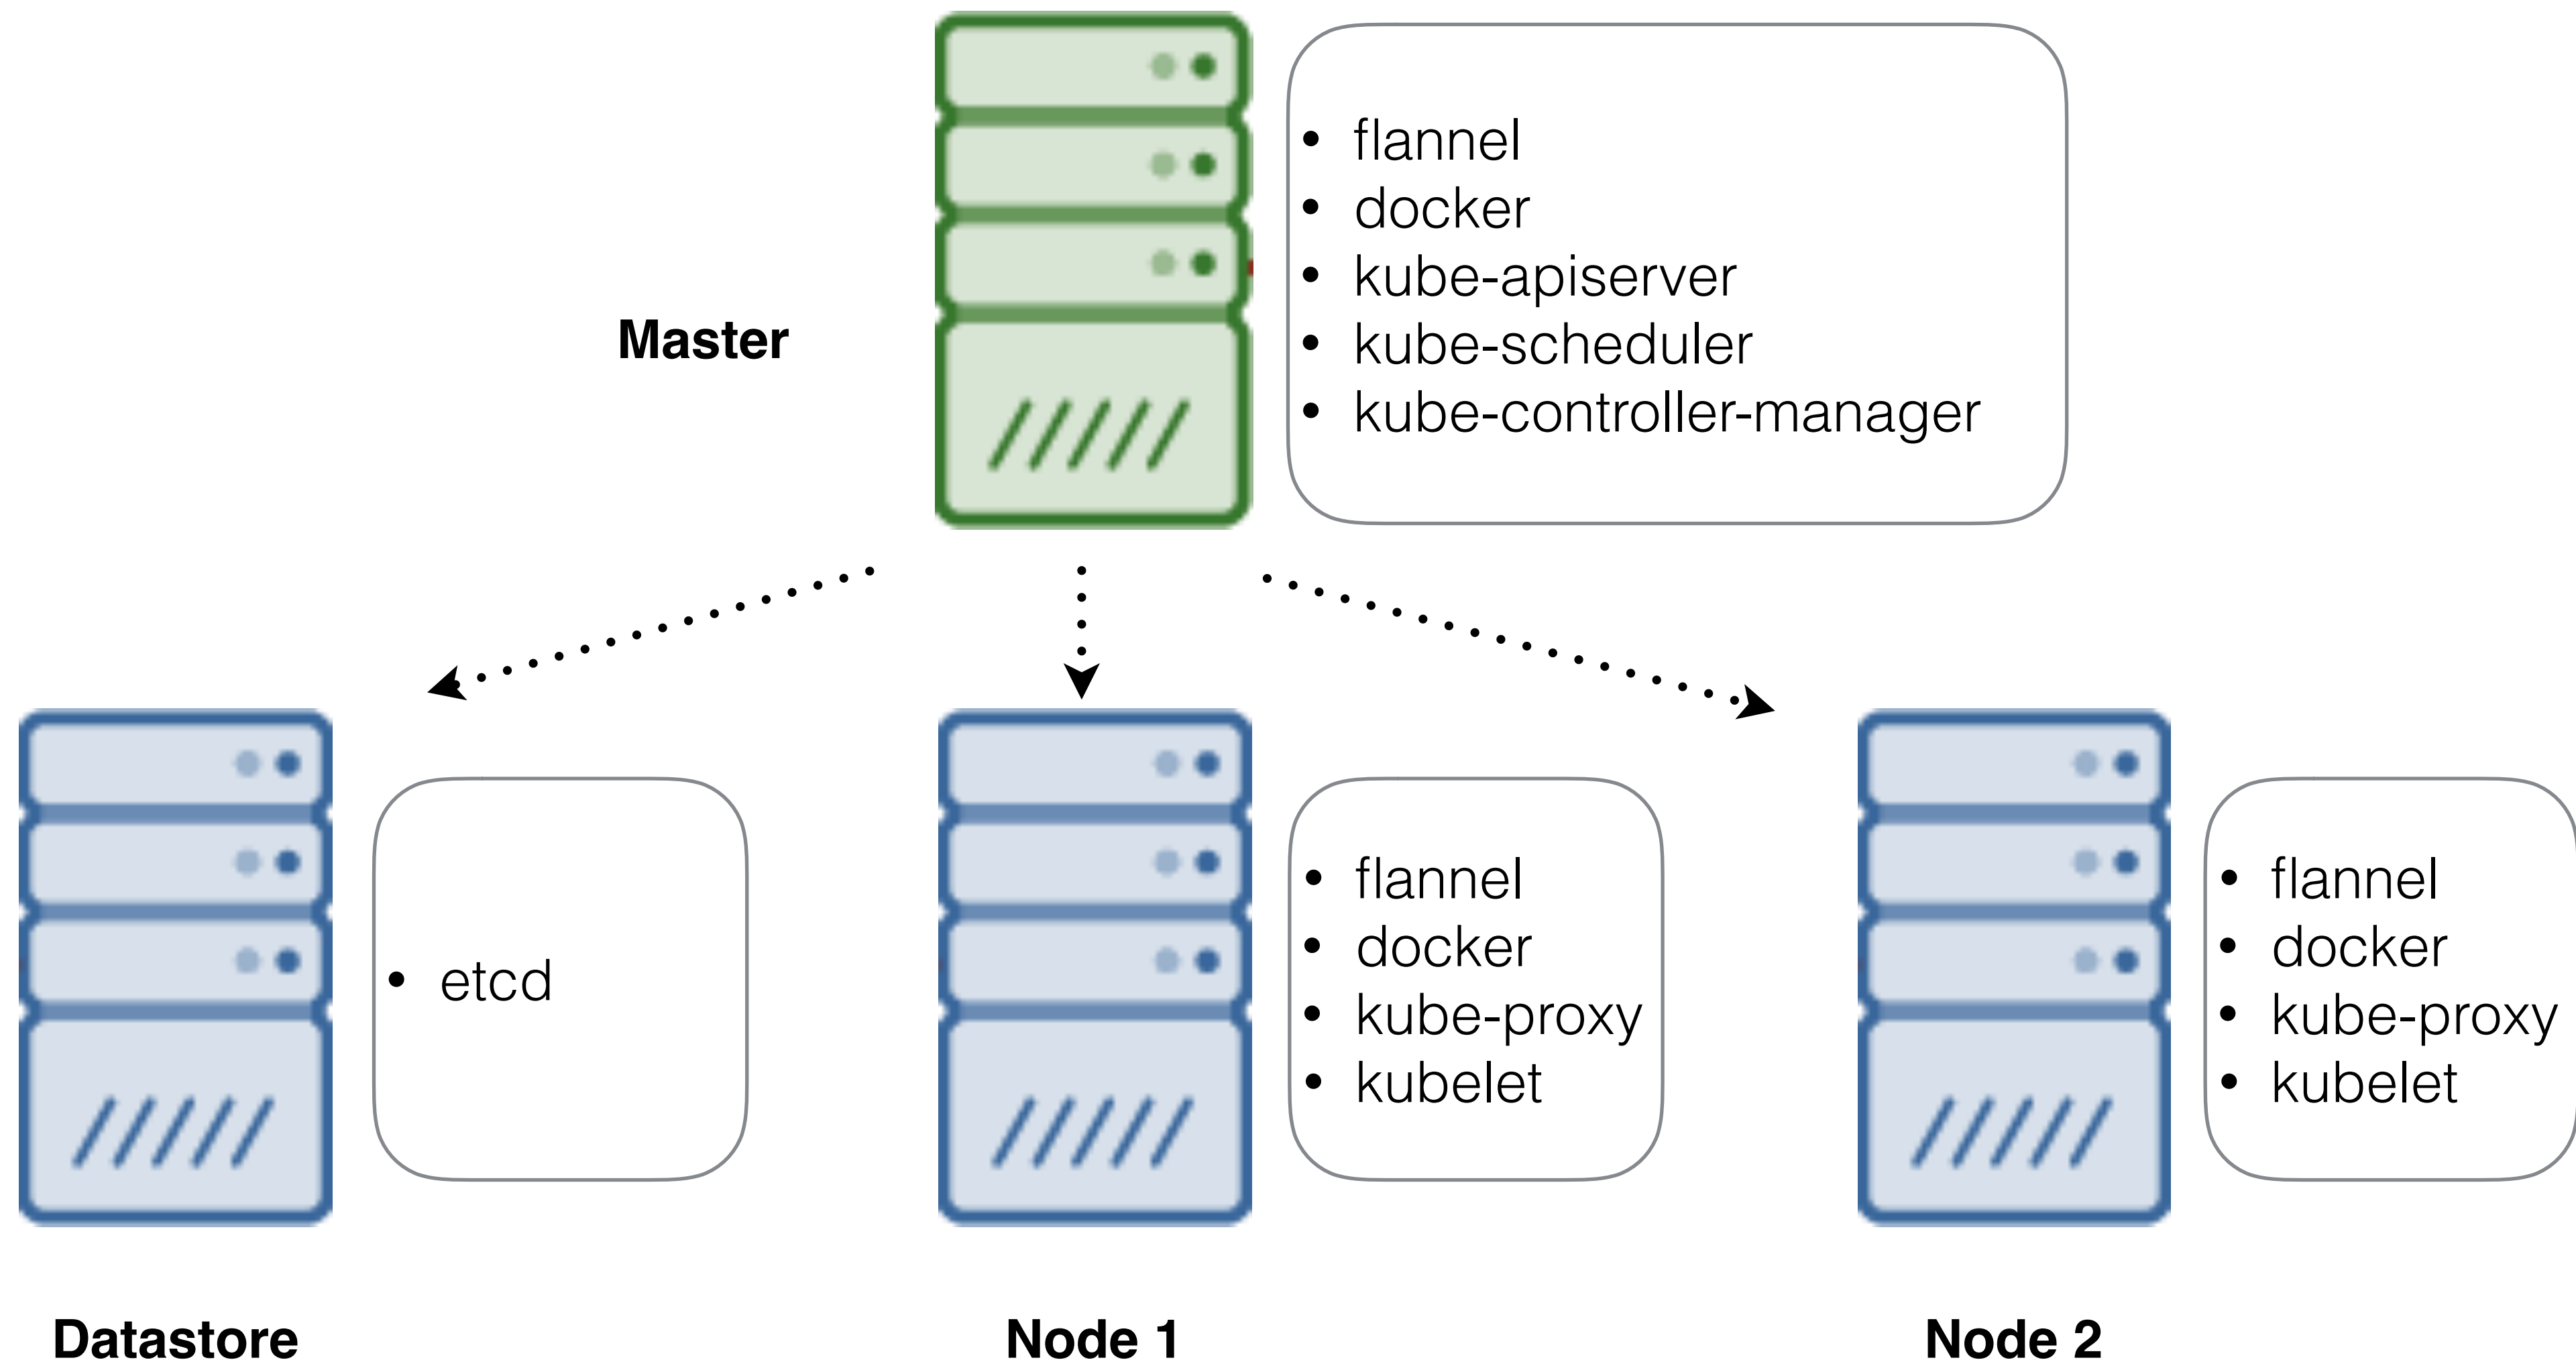


Welcome Kubernetes

- Kubernetes was conceived of and developed in a world where external developers were becoming interested in Linux containers.



Kubernetes inside



Kubernetes inside

- Kubernetes Master - centralized services that control the cluster (via the API, a controller manager and scheduler)
- Master Datastore - configuration storage which contains all persistent Kubernetes state (etcd).



Master


- kube-apiserver
- kube-scheduler
- kube-controller-manager



Datastore

- etcd

Kubernetes inside



- Nodes - VMs to place several pods and containers.
- Pods - a group of containers that must be placed on a single node and work together as a team.

Kubernetes inside

- Kubectl - cluster management console.
- Kubelet - as agent runs on every node and is responsible for driving Docker, reporting status to the master, and setting up node-level resources.
- Kube proxy - runs on every node and provides local containers a single network endpoint to reach an array of pods.



Master

- kubectl
- kube-proxy
- kubelet



Node

- kubectl
- kube-proxy
- kubelet

Kubernetes inside

- Labels - key/value tags attached to pods in order to organize a group for monitoring and management.
- Replication Controllers - agents that work to make sure that a horizontal scaling group or pod is reliably maintained.



Node 1

- name = webapp
- version = 0.1
- type = production




Node 2

- name = webapp
- version = 0.2
- type = stage


Kubernetes Cluster features

- Reliable container restart - monitor the health of a container and restart it when it fails.




Kubernetes Cluster features

- Self-healing - If a node fails can automatically reschedule work onto healthy nodes.



Kubernetes Cluster features

- High-cluster utilization
- can drive dramatically higher utilization compared to static manual placement.




Kubernetes Cluster features

- Rolling Update - ability to carry out the release without downtime and quickly return to the previous release in the event of failure.



Kubernetes Cluster features

- Horizontal scale and replication.
- Provides a flexible labeling system that allows both users and other systems to think in sets of containers.
- Automatic discovering.



Setup Kubernetes Cluster

```
# Clone k8sdemo repository
$ git clone https://github.com/takama/k8sdemo.git

# Setup ansible (DevOps automation/deployment tools)
$ yum install ansible

# Go to working directory
$ cd k8sdemo/ansible

# Setup environment
$ vim inventory/cluster
[master]
k8s-master.your-domain

[node]
k8s-node-01.your-domain
k8s-node-02.your-domain

[datastore:children]
master
```

```
# Run ansible playbook
$ ansible-playbook playbooks/cluster/setup.yml

PLAY [datastore] *****

TASK [setup] *****
ok: [k8s-master.openprovider.nl]

TASK [ping] *****
ok: [k8s-master.openprovider.nl]

PLAY [master] *****

TASK [setup] *****
ok: [k8s-master.openprovider.nl]

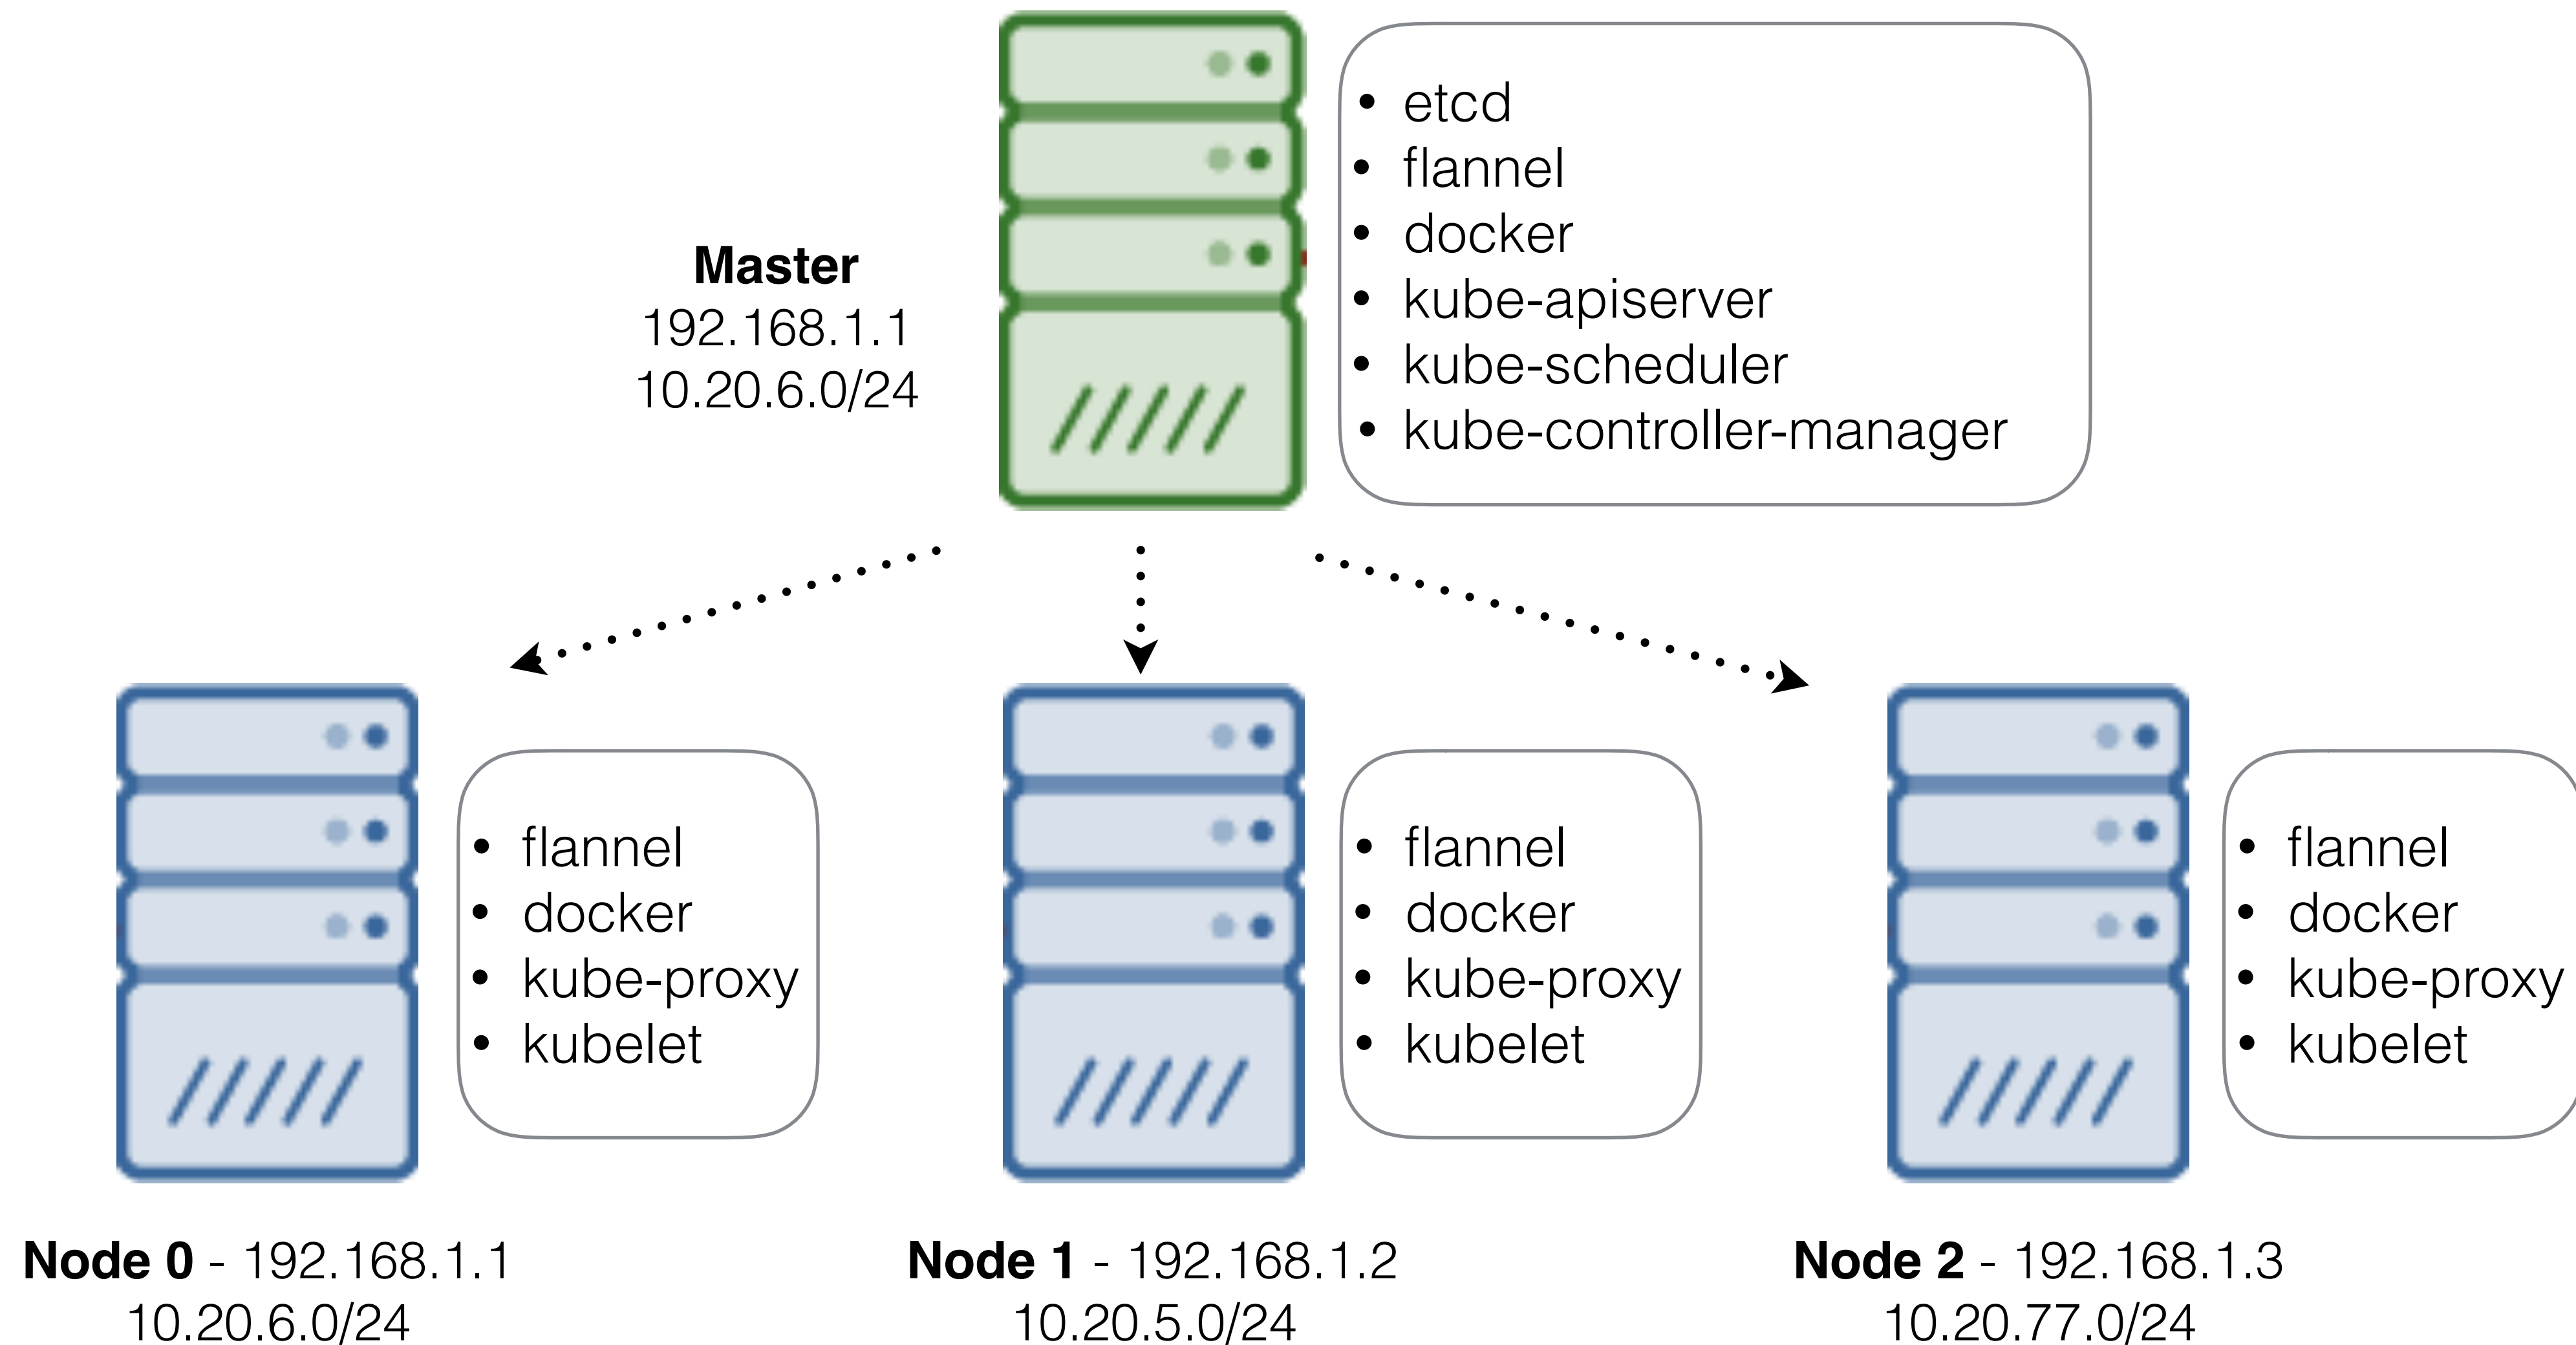
TASK [ping] *****
ok: [k8s-master.openprovider.nl]

PLAY [node] *****

TASK [setup] *****
ok: [k8s-node-01.openprovider.nl]
ok: [k8s-node-02.openprovider.nl]

TASK [ping] *****
ok: [k8s-node-01.openprovider.nl]
ok: [k8s-node-02.openprovider.nl]
```

Kubernetes Cluster



Kubernetes Dashboard

☰

kubernetes

Replication Controllers > webdb-controller

⋮

SCALE

EDIT

DELETE

Admin

Namespaces

Nodes

Persistent Volumes

Namespace

default ▼

Overview

Workloads

Deployments

Replica Sets

Replication Controllers

Daemon Sets

Pet Sets

Jobs

Pods

Services and discovery

Details

Status

Name: webdb-controller

Namespace: default

Label selector: name: webdb

Labels: name: webdb

Images: takama/webdb:0.1

Pods: 2 running

Services

Name	Labels	Cluster IP	Internal endpoints	External endpoints
✓ webdb	name: webdb	10.254.222.11	webdb:3306 TCP webdb:0 TCP	-

Pods

Name	Status	Restarts	Age	Cluster IP	CPU (cores)
✓ webdb-controller-1hedd	Running	0	8 hours	10.20.77.2	-
✓ webdb-controller-sgzs8	Running	0	8 hours	10.20.5.2	-

Events

Setup Kubernetes Dashboard

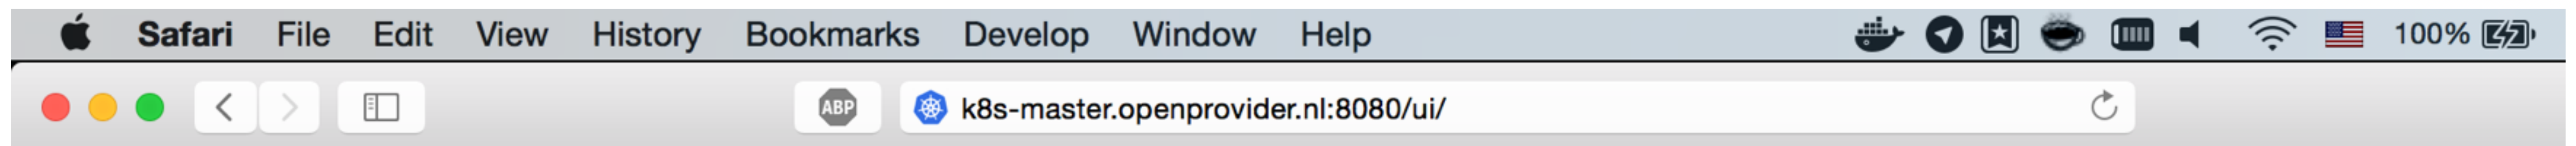
```
# Register kube-system namespace
$ kubectl create -f /etc/kubernetes/addons/kube-system-namespace.json
```

```
# Setup dashboard
$ kubectl create -f /etc/kubernetes/addons/dashboard.yaml
```

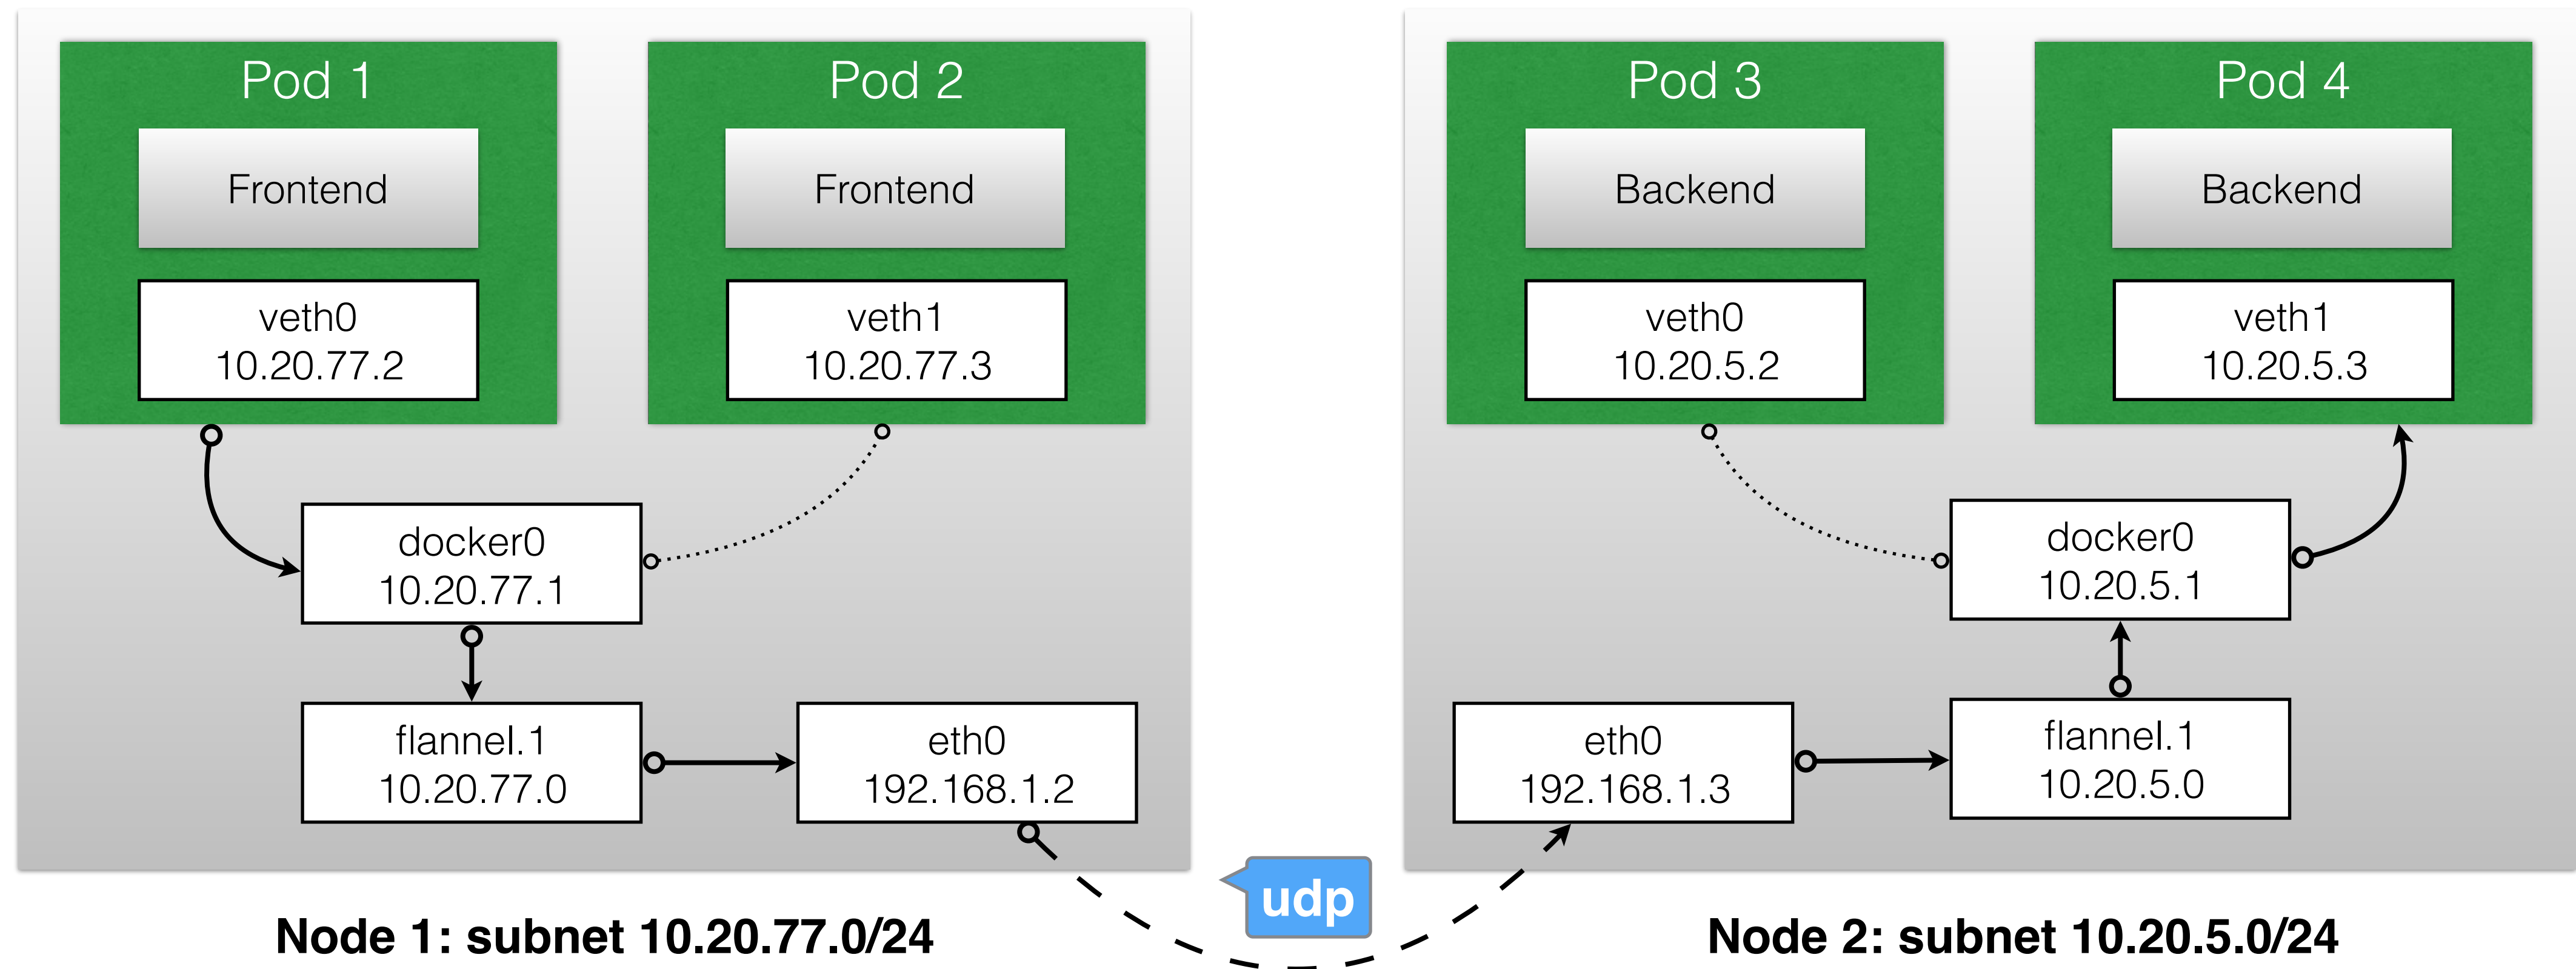
```
# Check installed dashboard
$ kubectl get svc --namespace=kube-system
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kube-dns	10.254.0.10	<none>	53/UDP,53/TCP	2d
kubernetes-dashboard	10.254.154.254	nodes	80/TCP	9d

```
# Open from browser
# http://master-host-name:8080/ui/
```



Overlay Network



Ping from Pod 1 to Pod 4: 10.2.77.2 -> 10.20.5.3

Setup Demo Manifests

```
# Run ansible playbook
$ ansible-playbook playbooks/cluster/demo.yml

PLAY [demo] *****

TASK [setup] *****
ok: [k8s-master.openprovider.nl]

TASK [ping] *****
ok: [k8s-master.openprovider.nl]

PLAY [demo] *****

TASK [setup] *****
ok: [k8s-master.openprovider.nl]

TASK [set_fact] *****
ok: [k8s-master.openprovider.nl]

TASK [demo : Check demo directory] *****
ok: [k8s-master.openprovider.nl]

TASK [demo : Setup demo service] *****
changed: [k8s-master.openprovider.nl] => (item=backend-svc.yaml)
changed: [k8s-master.openprovider.nl] => (item=frontend-svc.yaml)
changed: [k8s-master.openprovider.nl] => (item=backend-rc.yaml)
changed: [k8s-master.openprovider.nl] => (item=frontend-rc.yaml)
changed: [k8s-master.openprovider.nl] => (item=busybox.yaml)

PLAY RECAP *****
k8s-master.openprovider.nl : ok=6    changed=1    unreachable=0    failed=0
```

It copies manifests files into ~/services/

```
-rw-r--r--. 1 user 161 Sep  8 06:54 backend-svc.yaml
-rw-r--r--. 1 user 194 Sep  8 06:54 frontend-svc.yaml
-rw-r--r--. 1 user 342 Sep  8 06:54 backend-rc.yaml
-rw-r--r--. 1 user 356 Sep  8 06:55 frontend-rc.yaml
-rw-r--r--. 1 user 229 Sep  8 06:57 busybox.yaml
```


Overlay Network


```
# Prepare Pod manifest (optional if demo scripts installed)
$ cat > ~/services/busybox.yaml <<EOF
apiVersion: v1
kind: Pod
metadata:
  name: busybox
  namespace: default
spec:
  containers:
  - image: busybox
    command:
      - sleep
      - "3600"
    imagePullPolicy: IfNotPresent
    name: busybox
  restartPolicy: Always
EOF

# Create Pod
$ kubectl create -f ~/services/busybox.yaml
```

```
# ping Pod4 from Pod1 ( check corresponded IPs)
$ kubectl exec busybox -- ping 10.20.5.3 -s 1424
PING 10.20.5.3 (10.20.5.3): 1424 data bytes
1432 bytes from 10.20.5.3: seq=0 ttl=62 time=1.882 ms
1432 bytes from 10.20.5.3: seq=1 ttl=62 time=0.690 ms
1432 bytes from 10.20.5.3: seq=2 ttl=62 time=0.665 ms
1432 bytes from 10.20.5.3: seq=3 ttl=62 time=0.802 ms
1432 bytes from 10.20.5.3: seq=4 ttl=62 time=0.884 ms

# Show packets
$ tcpdump -i flannel.1 -nnA src 10.20.77.0
tcpdump: verbose output suppressed, use -v or -vv for
full protocol decode
listening on flannel.1, link-type EN10MB (Ethernet),
capture size 65535 bytes
10:53:52.184108 IP 10.20.77.0 > 10.20.5.3: ICMP echo
request, id 1536, seq 47, length 1424
E..... .?.&m
.M.
...../M..g.....
```

Demo Services



Prepare Demo Services

```
# Check component statuses
```

```
$ kubectl get cs
```

NAME	STATUS	MESSAGE	ERROR
controller-manager	Healthy	ok	
scheduler	Healthy	ok	
etcd-0	Healthy	{"health": "true"}	

```
# Register backend service
```

```
$ kubectl create -f ~/services/backend-svc.yaml
```

```
# Register frontend service
```

```
$ kubectl create -f ~/services/frontend-svc.yaml
```

```
# Register backend Replication Controller
```

```
$ kubectl create -f ~/services/backend-rc.yaml
```

```
# Register frontend Replication Controller
```

```
$ kubectl create -f ~/services/frontend-rc.yaml
```

```
# Check services
```

```
$ kubectl get svc
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	10.254.0.1	<none>	443/TCP	11d
webapp	10.254.101.185	192.168.1.1	3000/TCP	1m
webdb	10.254.105.176	<none>	3306/TCP	1m

```
# Check replication controllers
```

```
$ kubectl get rc
```

NAME	DESIRED	CURRENT	AGE
webapp-controller	2	2	1m
webdb-controller	2	2	1m

```
# Check pods
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kube-apiserver-k8s-master.openprovider.nl	1/1	Running	0	20h
kube-controller-manager-k8s-master.openprovider.nl	1/1	Running	0	20h
kube-scheduler-k8s-master.openprovider.nl	1/1	Running	0	20h
webapp-controller-isule	1/1	Running	0	1m
webapp-controller-j0emk	1/1	Running	0	1m
webdb-controller-hic12	1/1	Running	0	1m
webdb-controller-o33zd	1/1	Running	0	1m

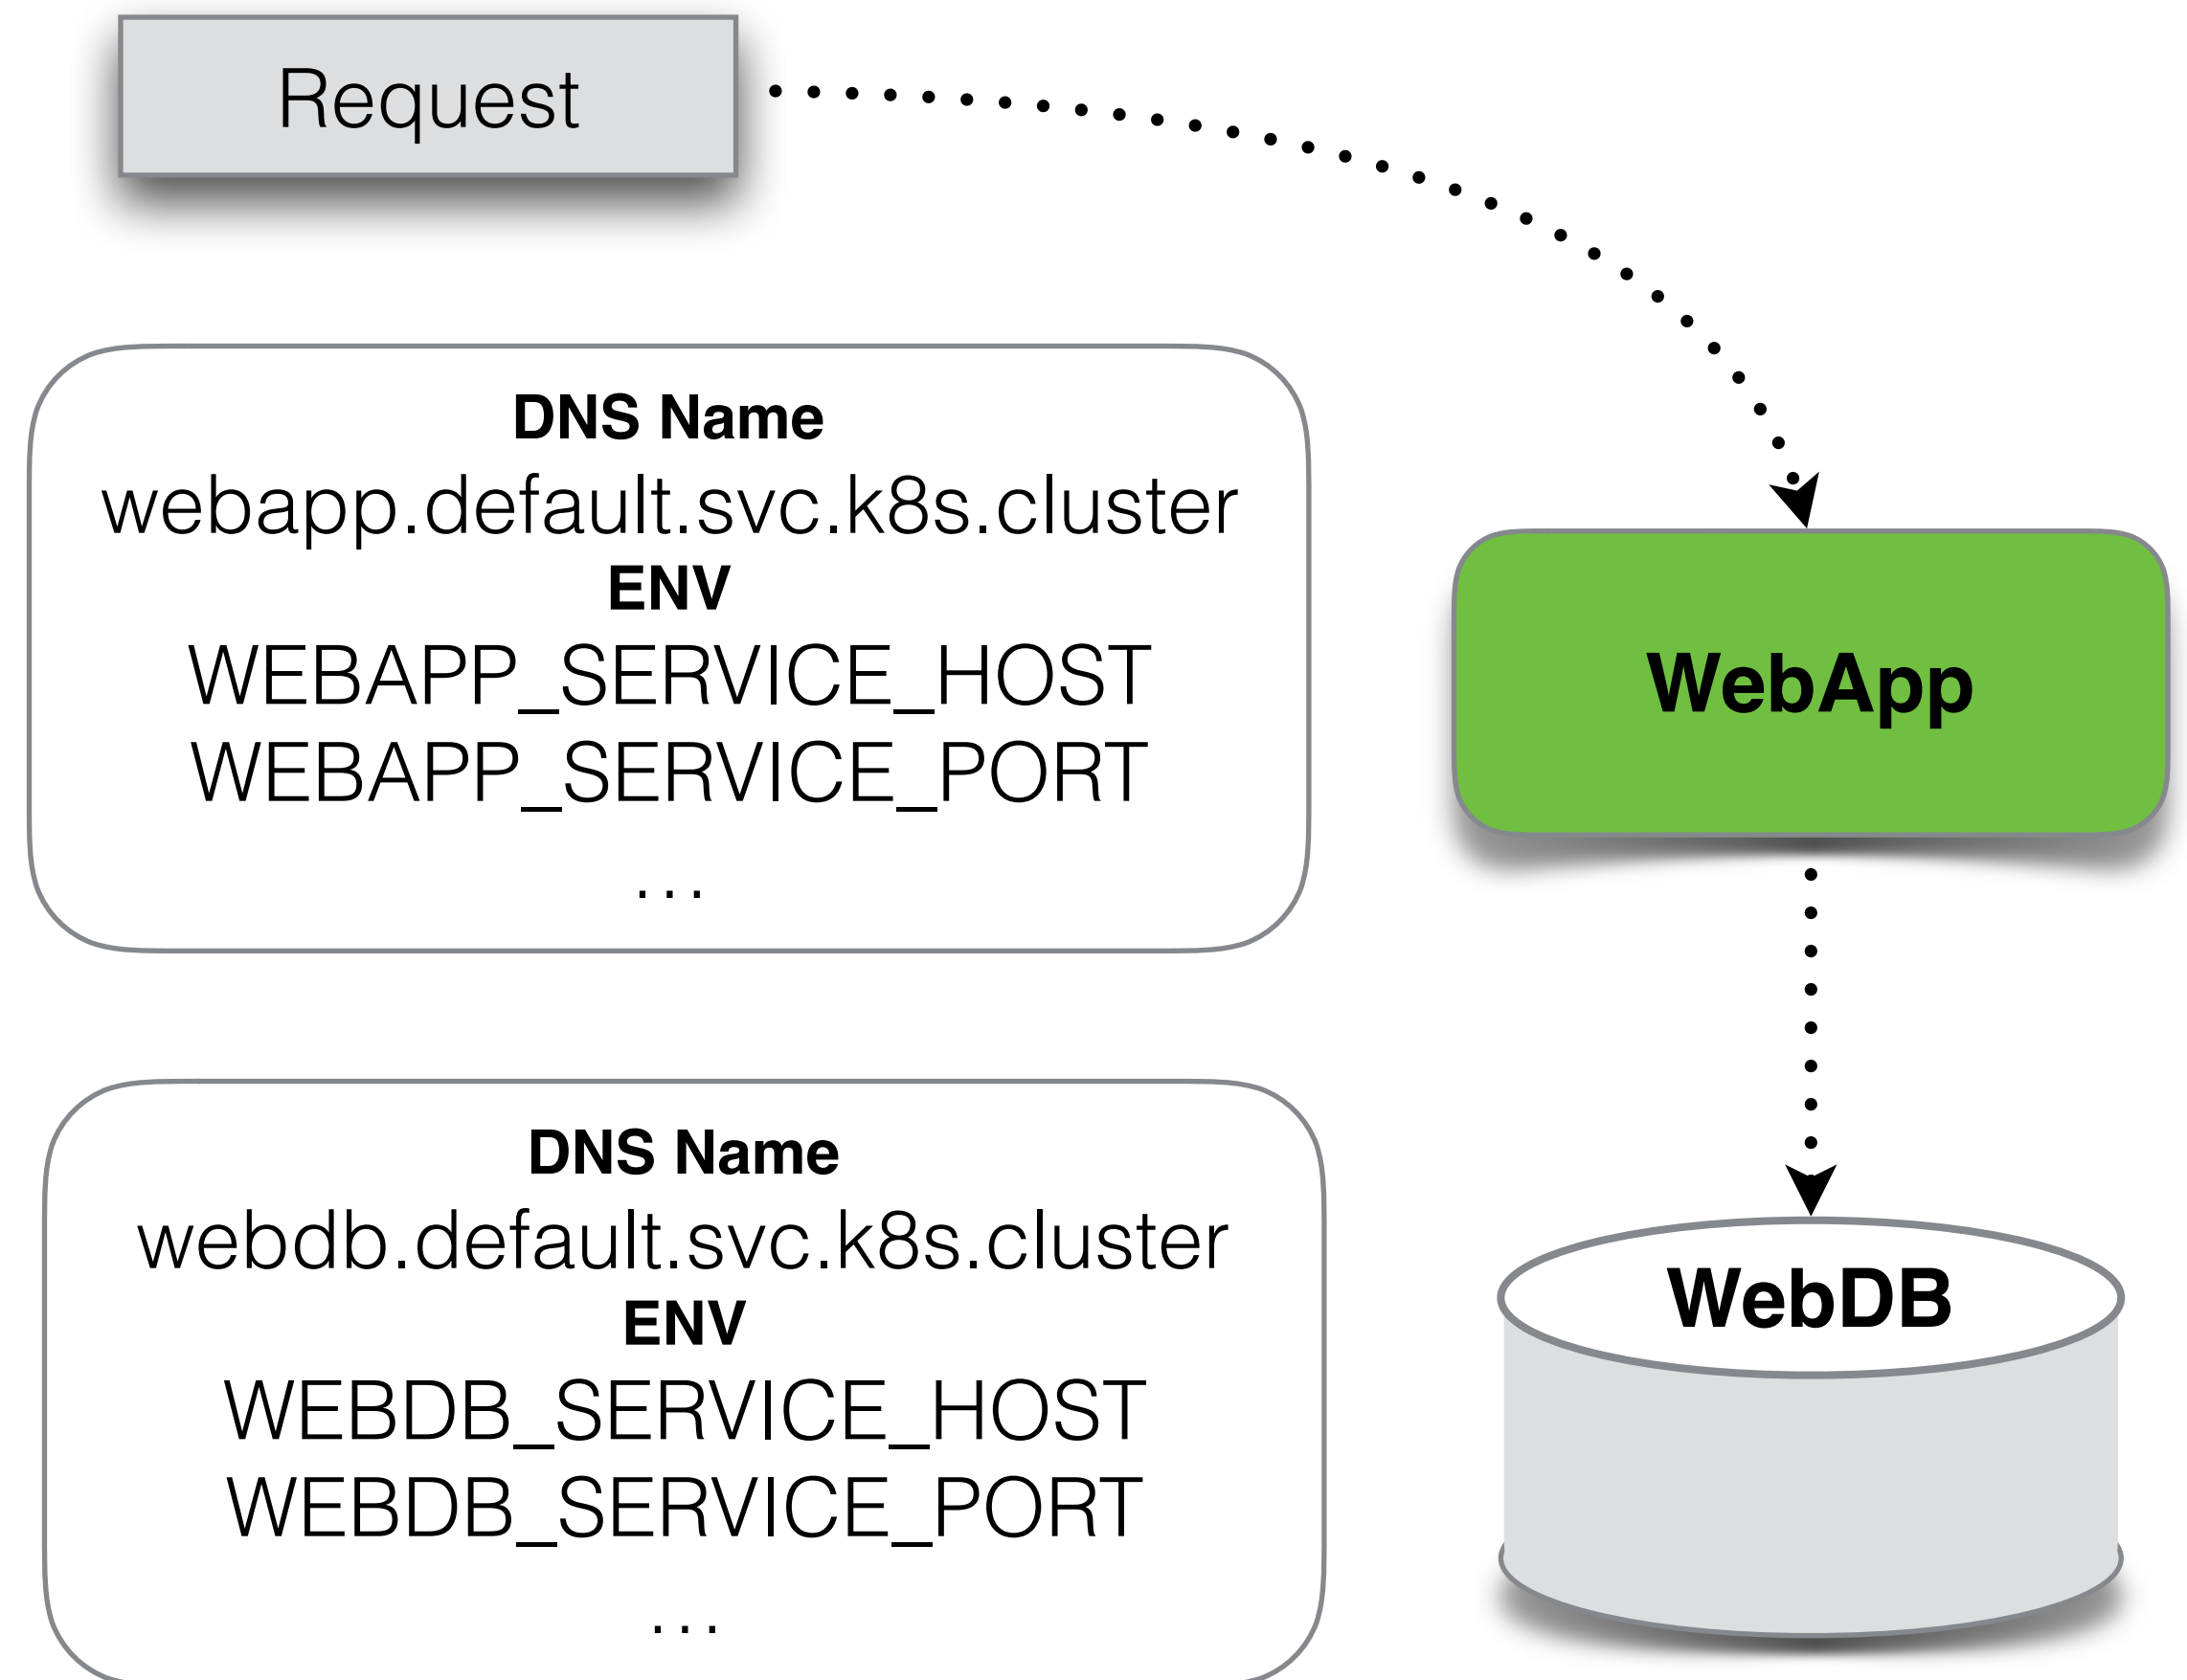
Use case 1: Relations

```
# Using Kube DNS for names resolving/discovering
#
# Template:
# < service name >< namespace >.svc.< cluster name >
#
# Lookup frontend service name from kubernetes
$ kubectl exec busybox -- nslookup webapp

Server:      10.254.0.10
Address 1: 10.254.0.10 kube-dns.kube-system.svc.k8s.cluster

Name:        webapp
Address 1: 10.254.159.24 webapp.default.svc.k8s.cluster

# Get backend container environments
$ docker inspect 2195619e1fdc | jq '.[0] | .Config.Env'
[
  "WEBAPP_SERVICE_HOST=10.254.159.24",
  "WEBDB_SERVICE_HOST=10.254.59.131",
  "WEBDB_SERVICE_PORT=3306",
  "WEBAPP_SERVICE_PORT=3000",
  "...",
  "GOSU_VERSION=1.7"
]
```



Use case 2: Scale

```
# Scale replicas from 2 to 3 for specified controller
$ kubectl scale --replicas=3 rc webapp-controller

replicationcontroller "webapp-controller" scaled

# Scale replicas from 3 to 2 for specified controller
$ kubectl scale --replicas=2 rc webapp-controller

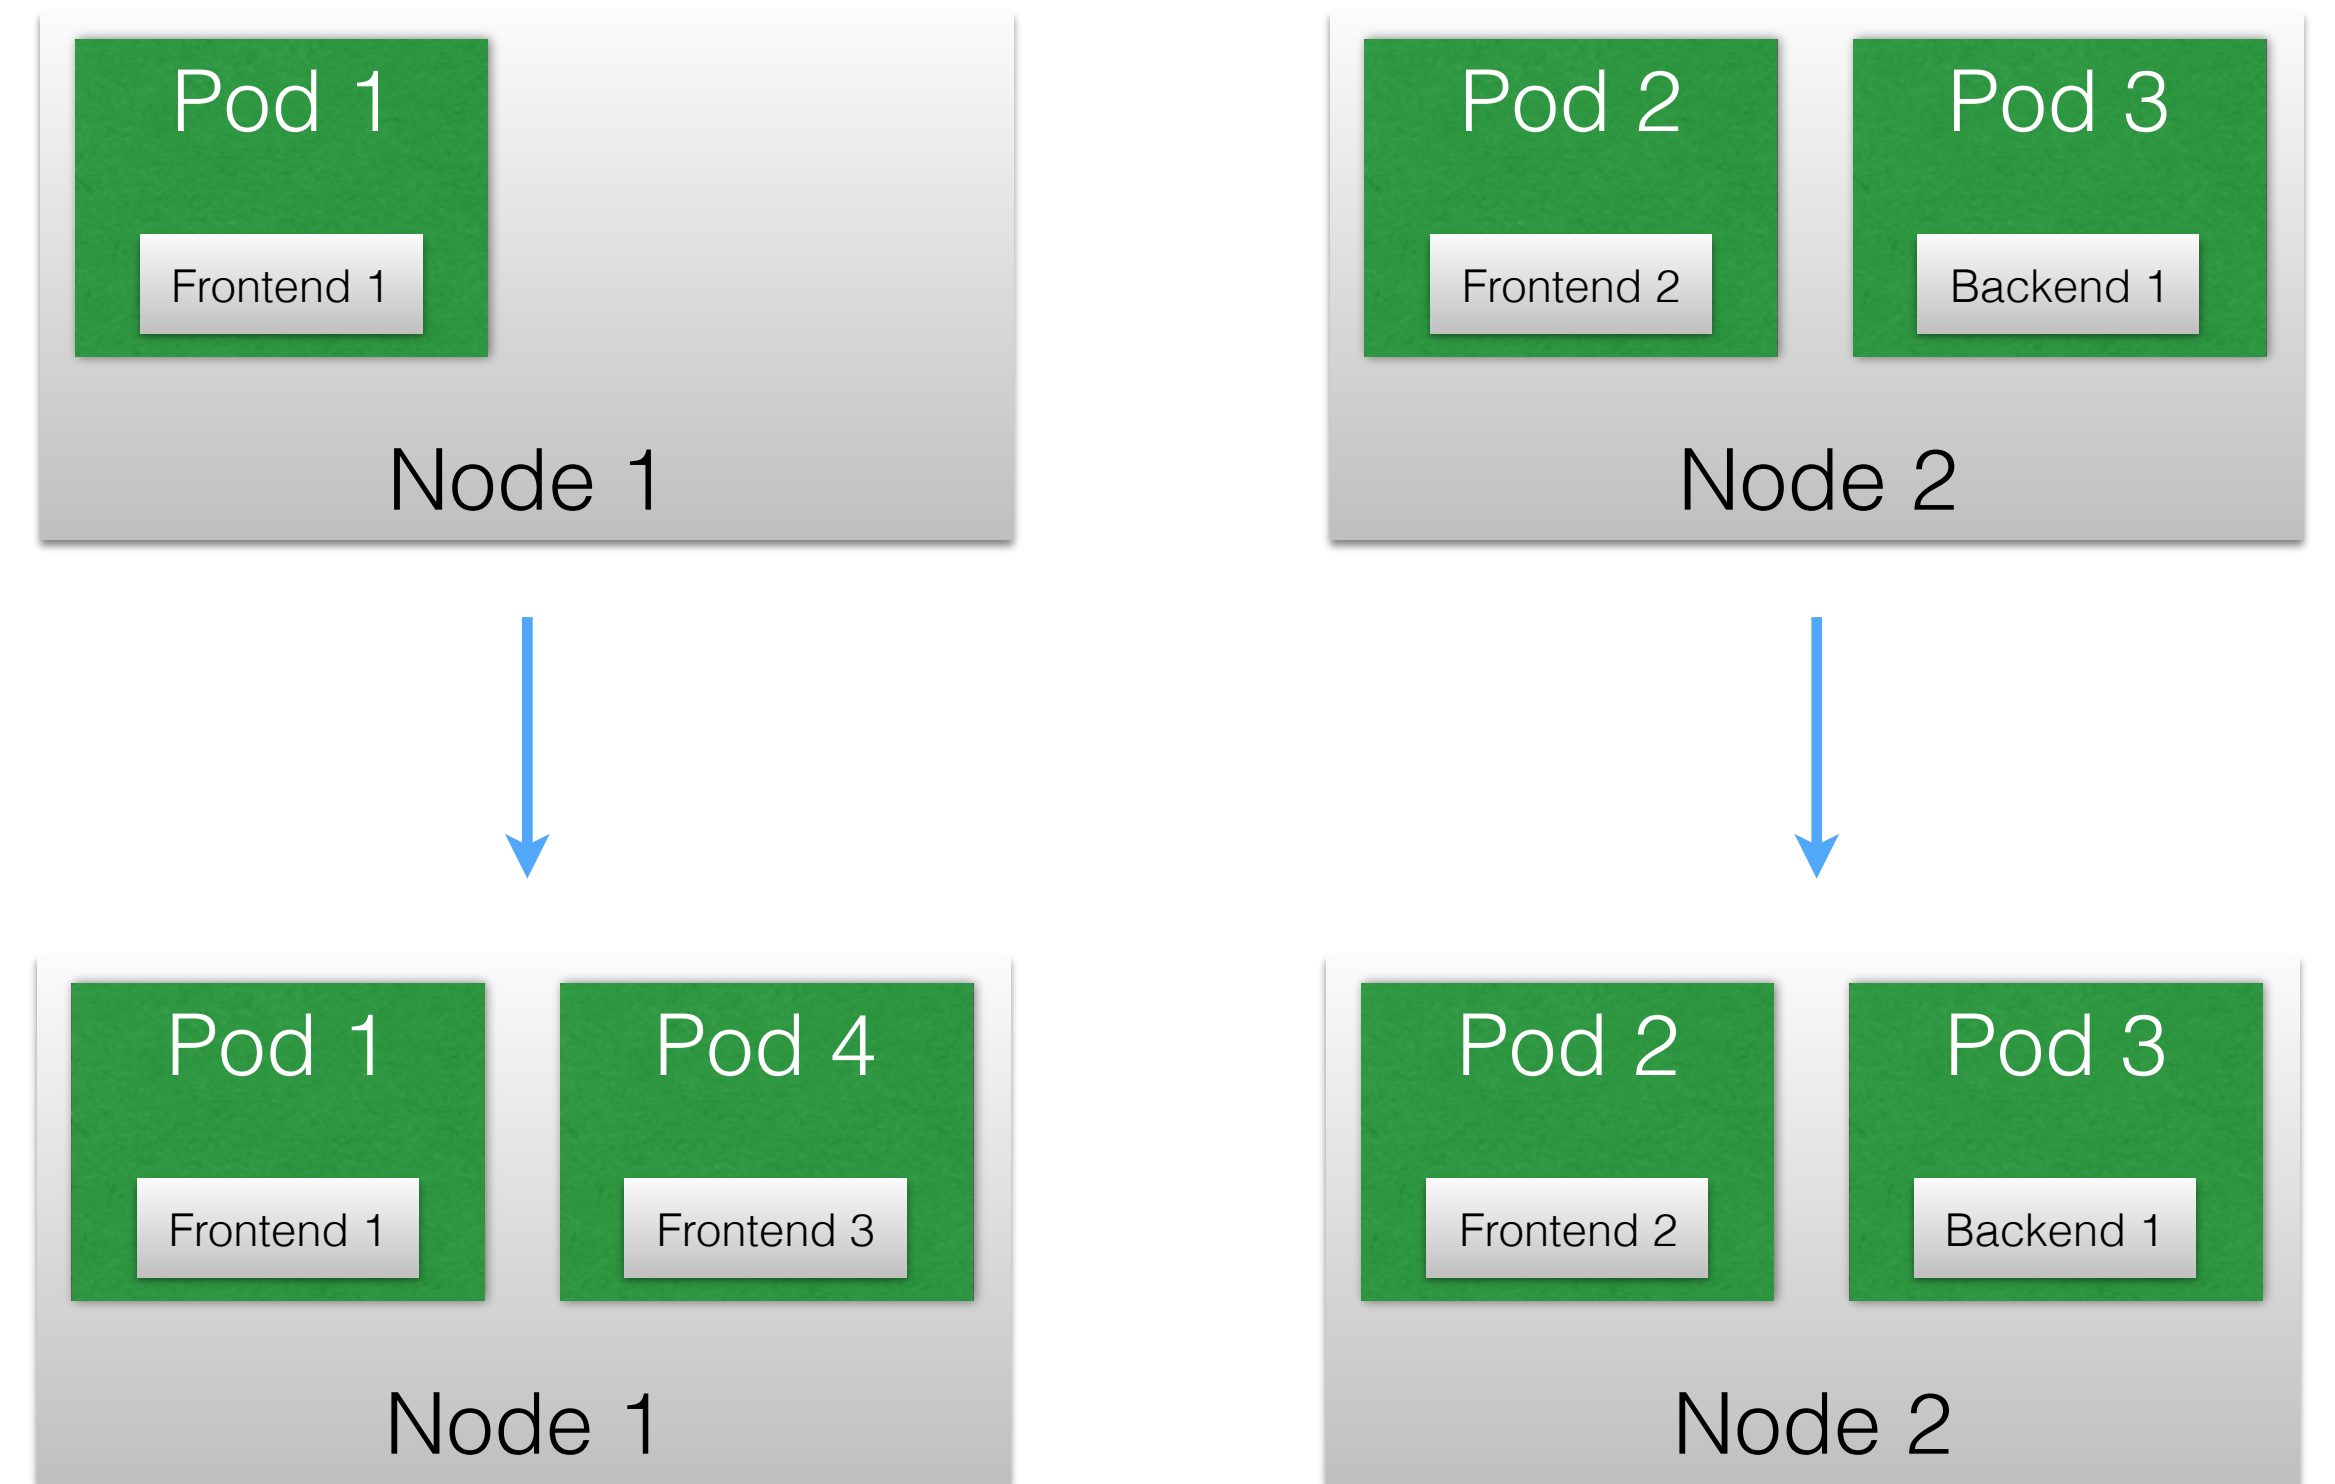
replicationcontroller "webapp-controller" scaled

# Scale replicas from 2 to 1 for specified controller
$ kubectl scale --replicas=1 rc webapp-controller

replicationcontroller "webapp-controller" scaled

# Scale replicas from 1 to 2 for specified controller
$ kubectl scale --replicas=2 rc webapp-controller

replicationcontroller "webapp-controller" scaled
```



Scale Frontend: 2 -> 3

Use case 3: Rolling Updates

```
# Rolling update to new image version
$ kubectl rolling-update webapp-controller --image=takama/webapp:0.2
```

```
Created webapp-controller-ed8873daf1dc43420f537d3585fdf337
Scaling up webapp-controller-ed8873daf1dc43420f537d3585fdf337 from 0 to 2, scaling down webapp-controller from 2 to 0 (keep 2 pods available, don't exceed 3 pods)
Scaling webapp-controller-ed8873daf1dc43420f537d3585fdf337 up to 1
Scaling webapp-controller down to 1
Scaling webapp-controller-ed8873daf1dc43420f537d3585fdf337 up to 2
Scaling webapp-controller down to 0
Update succeeded. Deleting old controller: webapp-controller
Renaming webapp-controller-ed8873daf1dc43420f537d3585fdf337 to webapp-controller
replicationcontroller "webapp-controller" rolling updated
```

```
# Rolling update to old image version
$ kubectl rolling-update webapp-controller --image=takama/webapp:0.1
```

```
Created webapp-controller-544345bd9584544cb36bc884bedb5829
Scaling up webapp-controller-544345bd9584544cb36bc884bedb5829 from 0 to 2, scaling down webapp-controller from 2 to 0 (keep 2 pods available, don't exceed 3 pods)
Scaling webapp-controller-544345bd9584544cb36bc884bedb5829 up to 1
Scaling webapp-controller down to 1
Scaling webapp-controller-544345bd9584544cb36bc884bedb5829 up to 2
Scaling webapp-controller down to 0
Update succeeded. Deleting old controller: webapp-controller
Renaming webapp-controller-544345bd9584544cb36bc884bedb5829 to webapp-controller
replicationcontroller "webapp-controller" rolling updated
```

Use case 3: Rolling Updates

```
# Rolling update to new image version with changing controller name
$ kubectl rolling-update webapp-controller webapp-controller-v2 --image=takama/webapp:0.2

# Rolling update to old image version with changing controller name
$ kubectl rolling-update webapp-controller-v2 webapp-controller --image=takama/webapp:0.1

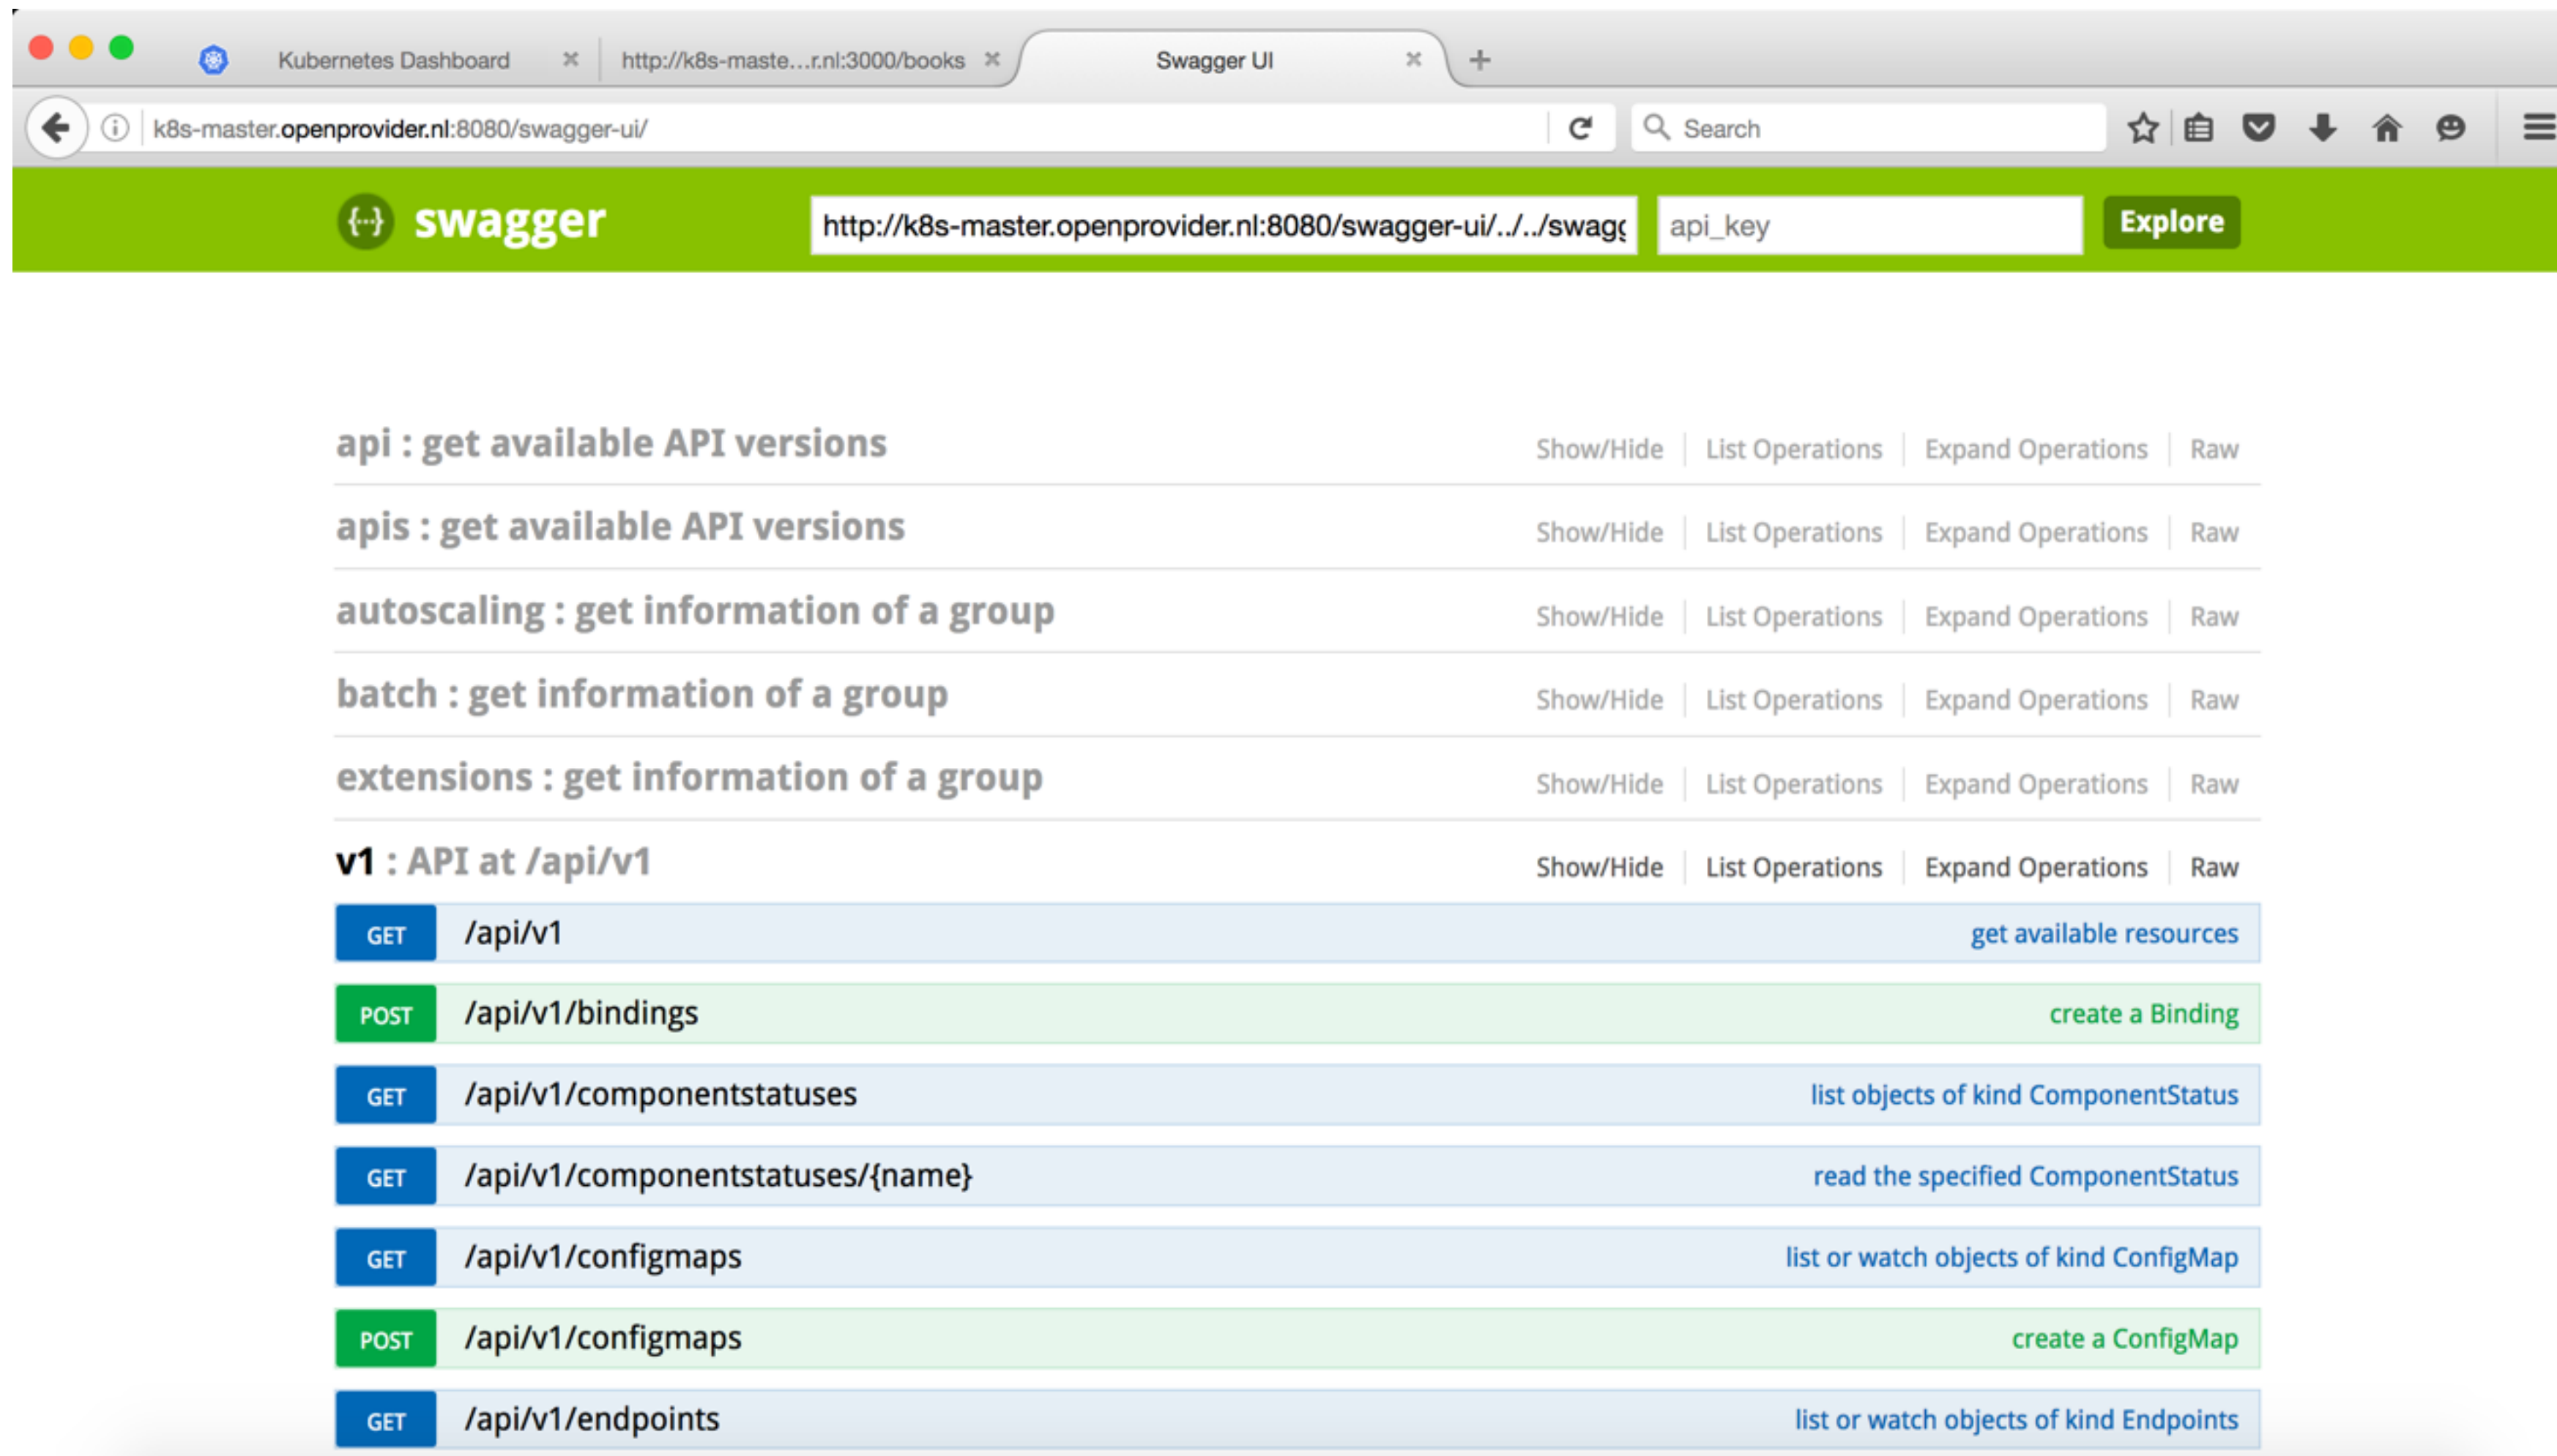
Created webapp-controller
Scaling up webapp-controller from 0 to 2, scaling down webapp-controller-v2 from 2 to 0 (keep 2 pods available, don't exceed 3 pods)
Scaling webapp-controller up to 1
Scaling webapp-controller-v2 down to 1
^C

# Rollback to previous state if something get wrong
$ kubectl rolling-update webapp-controller-v2 webapp-controller --rollback

Setting "webapp-controller-v2" replicas to 2
Continuing update with existing controller webapp-controller-v2.
Scaling up webapp-controller-v2 from 1 to 2, scaling down webapp-controller from 1 to 0 (keep 2 pods available, don't exceed 3 pods)
Scaling webapp-controller-v2 up to 2
Scaling webapp-controller down to 0
Update succeeded. Deleting webapp-controller

# Repeat rolling update to old image version with changing controller name
$ kubectl rolling-update webapp-controller-v2 webapp-controller --image=takama/webapp:0.1
```

http://<KUBE_MASTER_IP>:8080/swagger-ui/



Thank you

- Sources: <https://github.com/takama/k8sdemo>
- Docs: <https://access.redhat.com/documentation/en/red-hat-enterprise-linux-atomic-host/version-7/getting-started-with-containers/>